LLM optimization techniques

LLM optimization techniques

Distillation

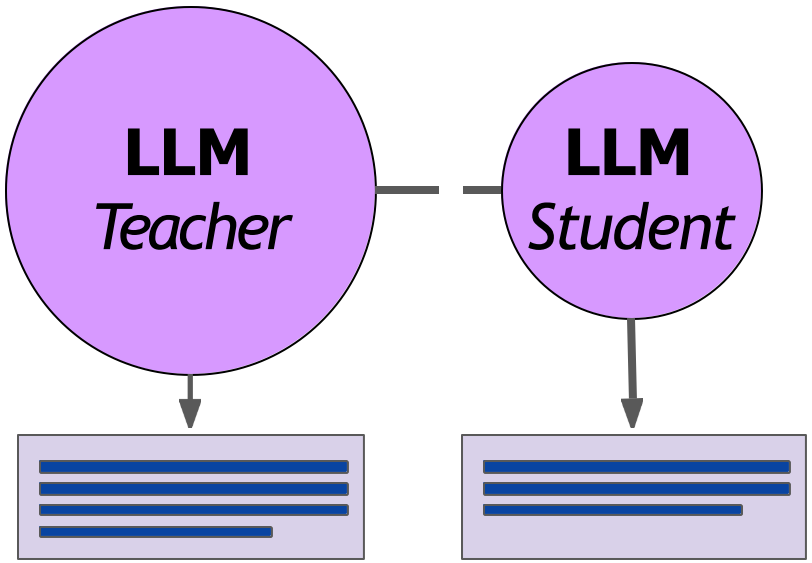

Quantization

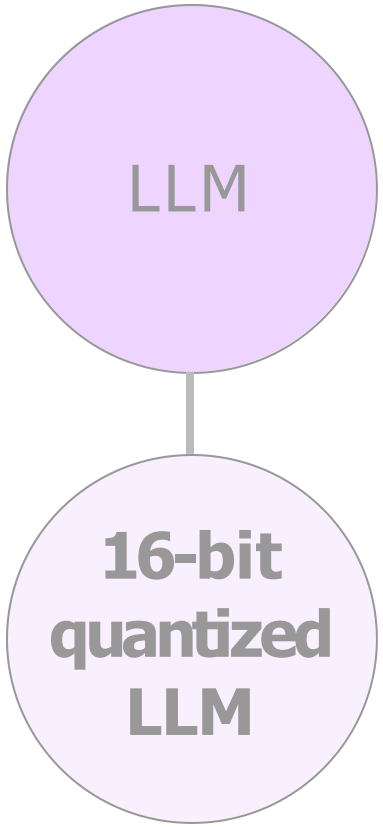

Pruning

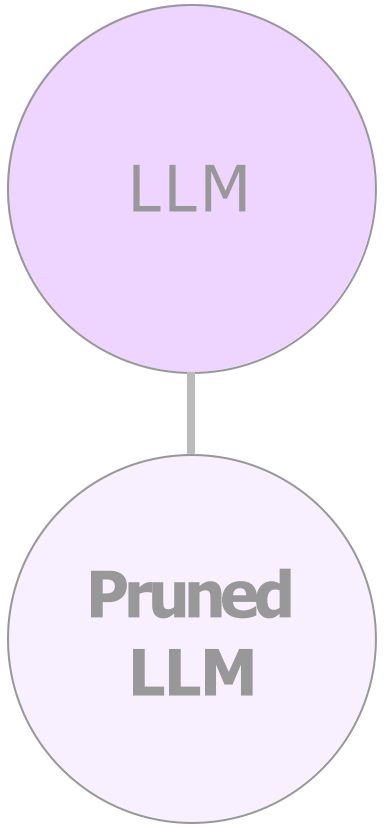

LLM optimization techniques

Distillation

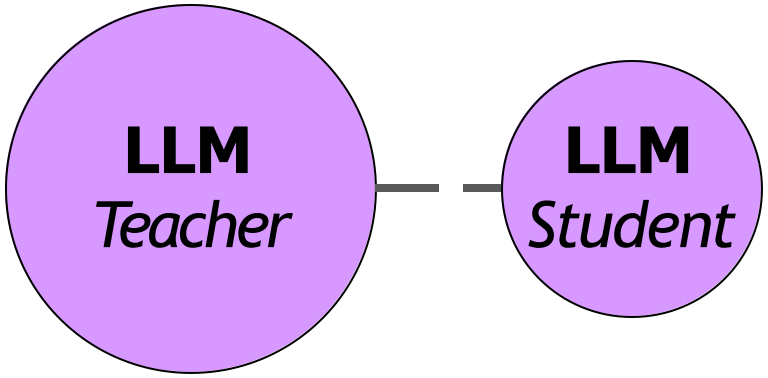

Quantization

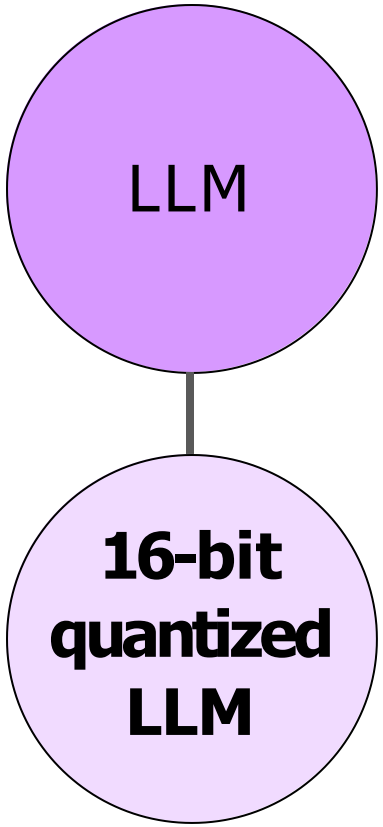

Pruning

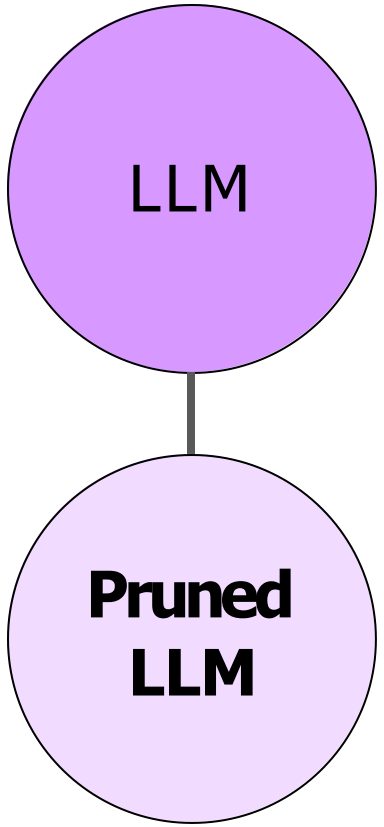

Distillation

Train a smaller student model from a larger teacher model

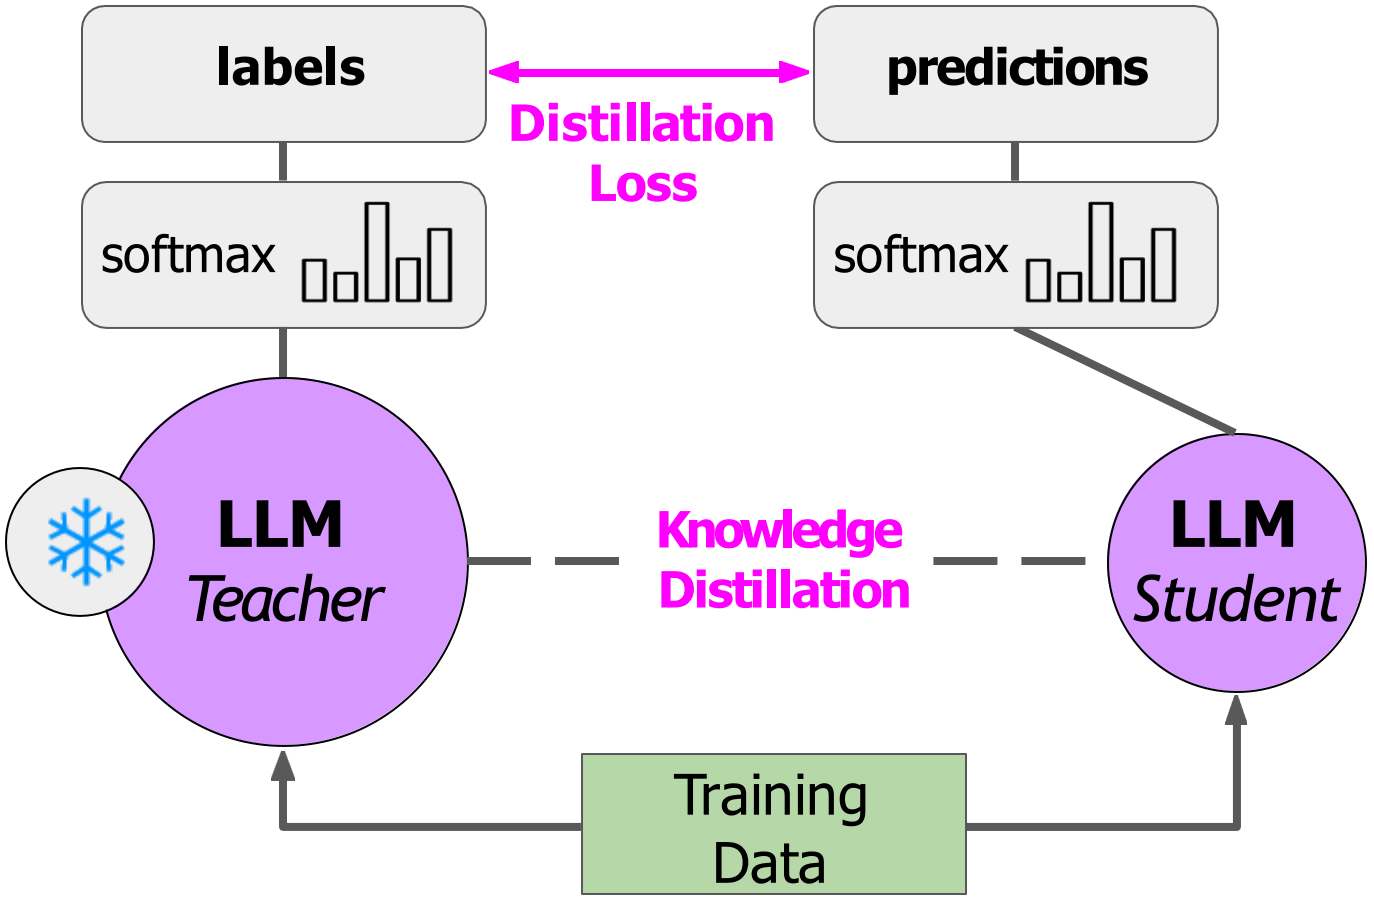

Distillation

Train a smaller student model from a larger teacher model

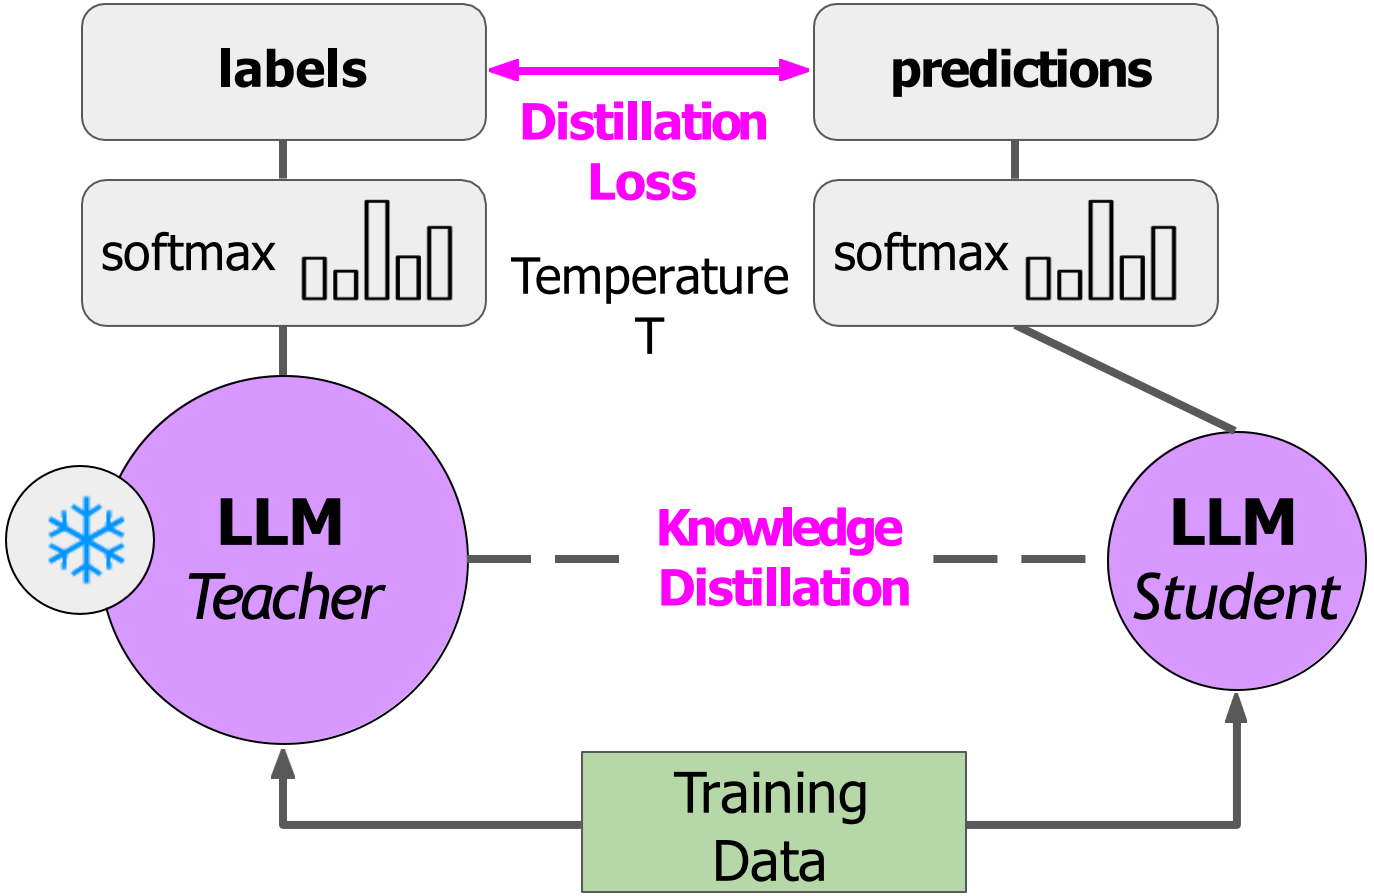

Distillation

Train a smaller student model from a larger teacher model

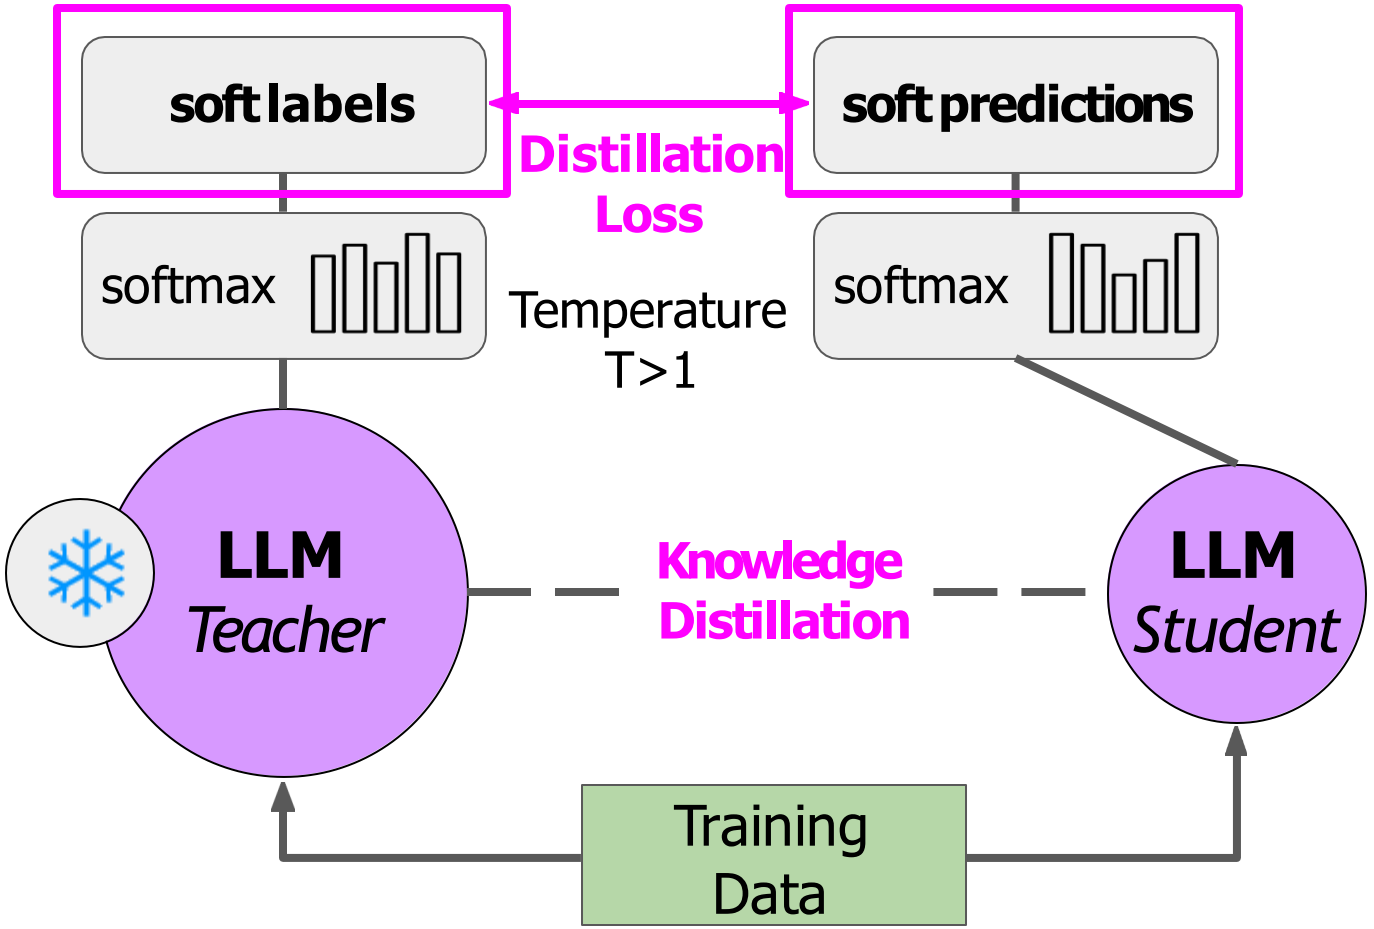

Distillation

Train a smaller student model from a larger teacher model

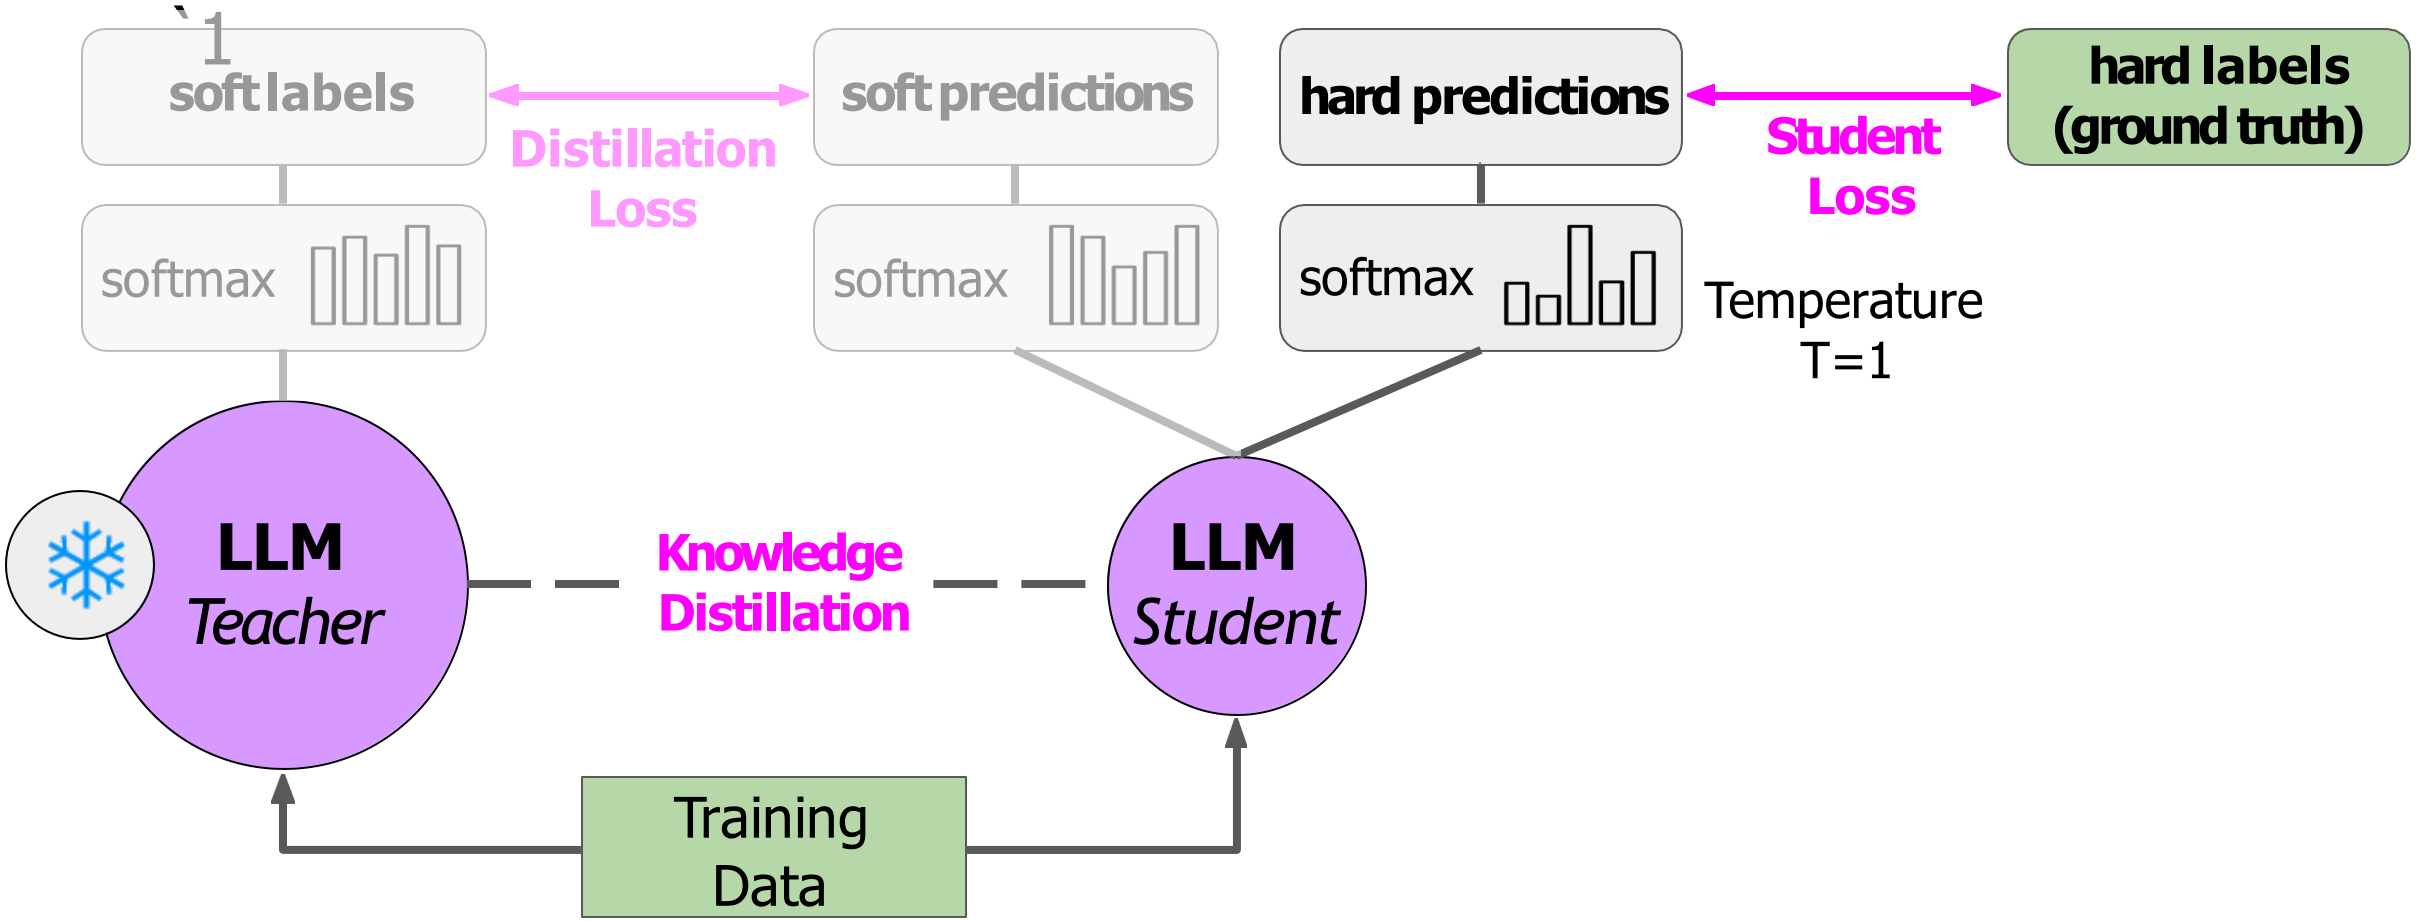

Distillation

Train a smaller student model from a larger teacher model

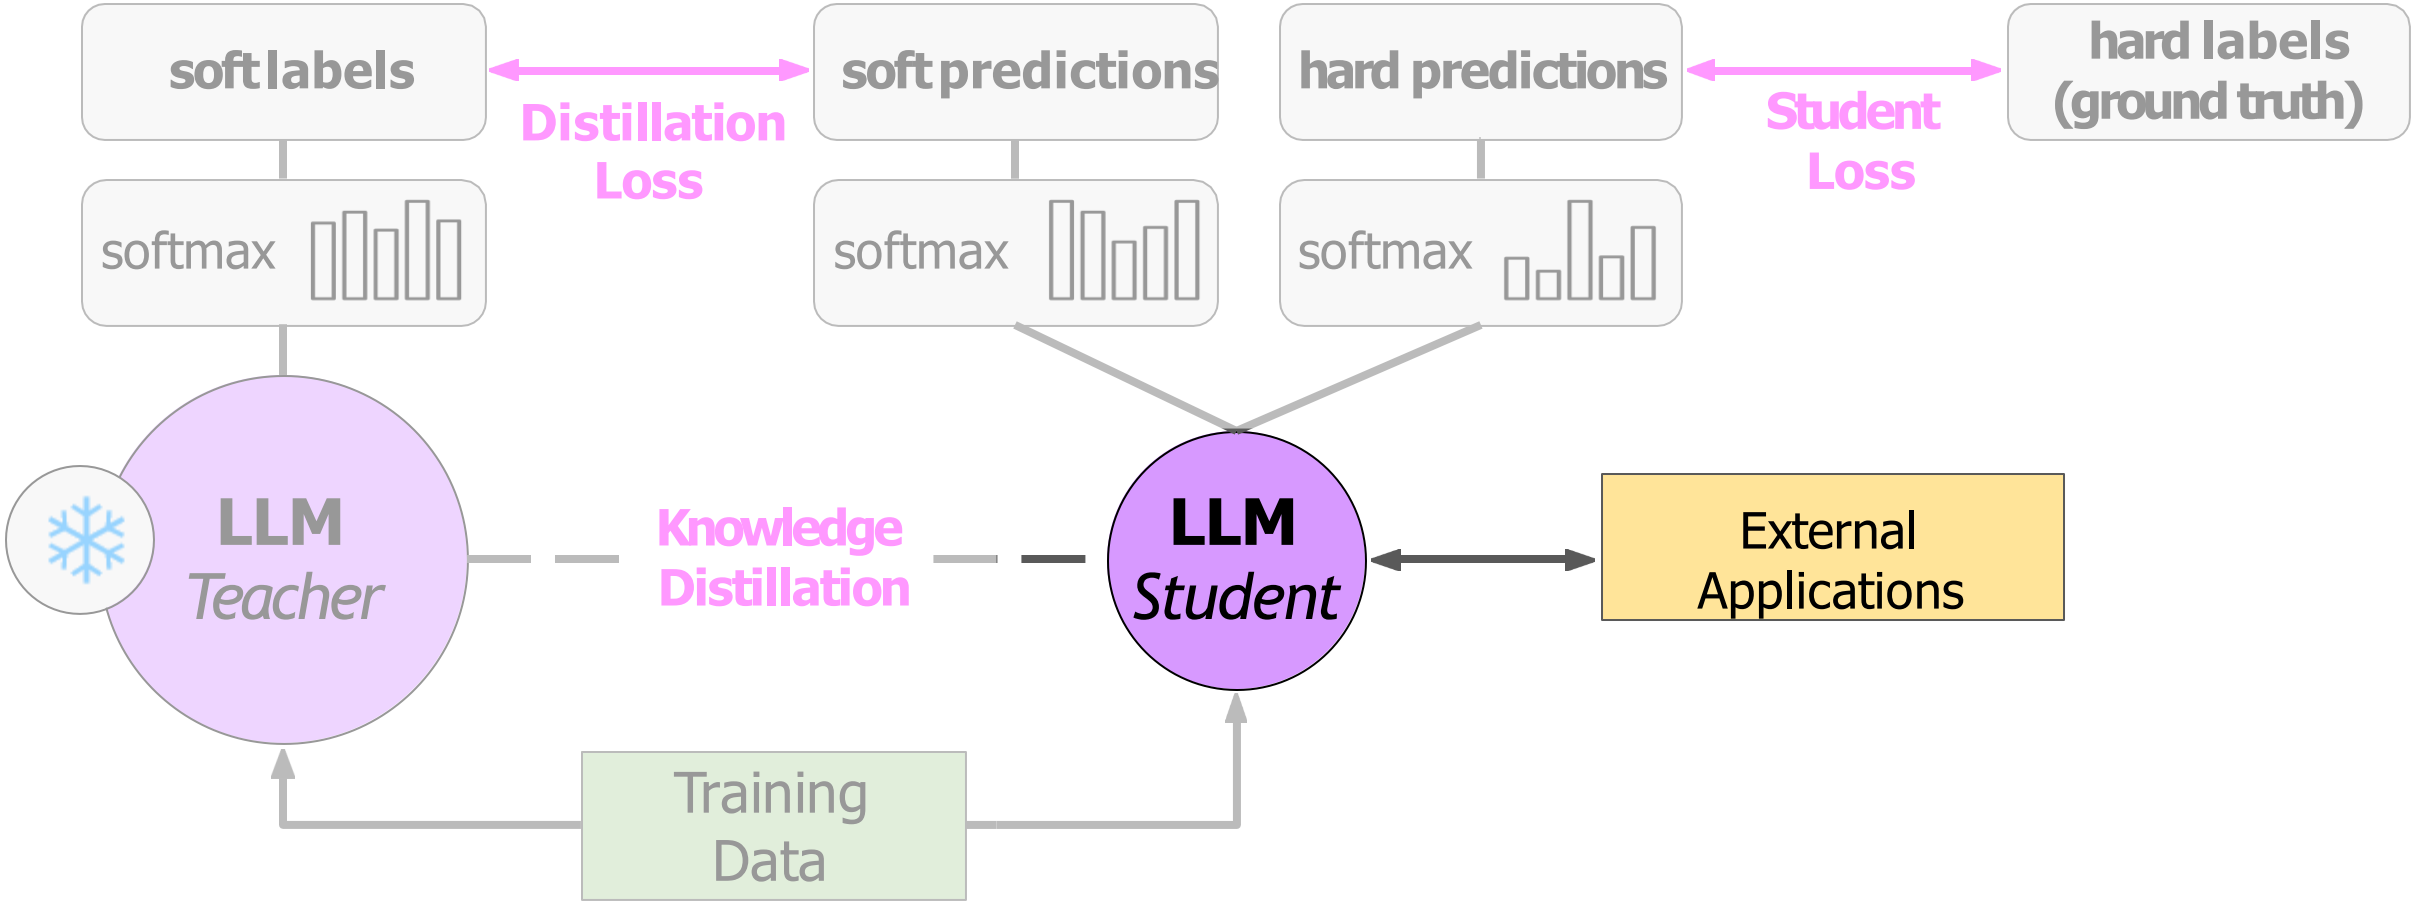

Backprogation

Backprogation update weights of LLM Students